

EXPERIMENT 1

Aim: Introduction to Data science and Data preparation using Pandas steps.

- Load data in Pandas.
- Description of the dataset.
- Drop columns that aren't useful.
- Drop rows with maximum missing values.
- Take care of missing data.
- Create dummy variables.
- Find out outliers (manually)
- Standardization and normalization of columns

Dataset: Financial Risk Assessment

Link: <https://www.kaggle.com/datasets/preethamgouda/financial-risk>

Steps:

1) Loading data in Pandas and extracting information about the dataset.

To load a file onto python for analysis, we need to make use of the pandas library. It gives us functionalities to read a CSV (Comma Separated Values) file and perform various functions on it.

Commands: `import pandas as pd` (Importing the pandas library onto Google Colab Notebook) `df = pd.read_csv()` (Mounts and reads the file in Python and assigns it to variable df for ease of use further)

(Note: Replace with the actual path of the file in “”)

`dataset.info()`: This command gives all the information about the features (columns) of the dataset and the data type of each of these columns. It also gives a summary of all the values in the dataset.

```

import pandas as pd
import numpy as np

# loading the dataset to pandas df
dataset = pd.read_csv("/content/financial_risk_assessment.csv")
dataset.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15000 entries, 0 to 14999
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   Age                   15000 non-null  int64  
 1   Gender                15000 non-null  object  
 2   Education Level       15000 non-null  object  
 3   Marital Status        15000 non-null  object  
 4   Income                12750 non-null  float64 
 5   Credit Score          12750 non-null  float64 
 6   Loan Amount           12750 non-null  float64 
 7   Loan Purpose          15000 non-null  object  
 8   Employment Status     15000 non-null  object  
 9   Years at Current Job  15000 non-null  int64  
10   Payment History       15000 non-null  object  
11   Debt-to-Income Ratio  15000 non-null  float64 
12   Assets Value          12750 non-null  float64 
13   Number of Dependents  12750 non-null  float64 
14   City                  15000 non-null  object  
15   State                 15000 non-null  object  
16   Country               15000 non-null  object  
17   Previous Defaults      12750 non-null  float64 
18   Marital Status Change 15000 non-null  int64  
19   Risk Rating           15000 non-null  object  
dtypes: float64(7), int64(3), object(10)
memory usage: 2.3+ MB

```

- 2) `df.head()`: As mentioned before, head function give us the first 5 rows of the dataset. This allows for the user to get an overview on what values are being listed in the dataset.

```
[ ] dataset.head()
```

	Age	Gender	Education Level	Marital Status	Income	Credit Score	Loan Amount	Loan Purpose	Employment Status	Years at Current Job	Payment History	Debt-to-Income Ratio	Assets Value	Number of Dependents	City	State	Country	Previous Defaults	Marital Status Change
0	49	Male	PhD	Divorced	72799.0	688.0	45713.0	Business	Unemployed	19	Poor	0.154313	120228.0	0.0	Port Elizabeth	AS	Cyprus	2.0	2
1	57	Female	Bachelor's	Widowed	NaN	690.0	33835.0	Auto	Employed	6	Fair	0.148920	55849.0	0.0	North Catherine	OH	Turkmenistan	3.0	2
2	21	Non-binary	Master's	Single	55687.0	600.0	36623.0	Home	Employed	8	Fair	0.362398	180700.0	3.0	South Scott	OK	Luxembourg	3.0	2
3	59	Male	Bachelor's	Single	26508.0	622.0	26541.0	Personal	Unemployed	2	Excellent	0.454964	157319.0	3.0	Robinhaven	PR	Uganda	4.0	2
4	25	Non-binary	Bachelor's	Widowed	49427.0	766.0	36528.0	Personal	Unemployed	10	Fair	0.143242	287140.0	NaN	New Heather	IL	Namibia	3.0	1

- 3) `dataset.shape()`: returns the dimensions of the dataset as a tuple (rows, columns), helping to understand its size.

```
[ ] dataset.shape
```

```
(15000, 20)
```

- 4) Describe the dataset

`dataset.describe()`: provides statistical summaries of numerical columns, including count, mean, standard deviation, min, max, and quartiles (25%, 50%, 75%).

dataset.describe()

	Age	Income	Credit Score	Loan Amount	Years at Current Job	Debt-to-Income Ratio	Assets Value	Number of Dependents	Previous Defaults	Marital Status Change
count	15000.000000	12750.000000	12750.000000	12750.000000	15000.000000	15000.000000	12750.000000	12750.000000	12750.000000	15000.000000
mean	43.452667	69933.398510	699.109098	27450.010902	9.476267	0.350438	159741.497176	2.02651	1.992471	0.998467
std	14.910732	29163.626207	57.229465	12949.940135	5.769707	0.143919	80298.115832	1.41130	1.416909	0.813782
min	18.000000	20005.000000	600.000000	5000.000000	0.000000	0.100004	20055.000000	0.000000	0.000000	0.000000
25%	31.000000	44281.500000	650.000000	16352.500000	4.000000	0.227386	90635.250000	1.000000	1.000000	0.000000
50%	43.000000	69773.000000	699.000000	27544.000000	9.000000	0.350754	159362.000000	2.000000	2.000000	1.000000
75%	56.000000	95922.750000	748.000000	38547.500000	15.000000	0.476095	228707.000000	3.000000	3.000000	2.000000
max	69.000000	119997.000000	799.000000	49998.000000	19.000000	0.599970	299999.000000	4.000000	4.000000	2.000000

If the parameter of include="all" is included { df.describe(include="all") }, this includes even the non numeric values and gives some more information on fields such as count of unique values, top value, etc.

dataset.describe(include="all")

	Age	Gender	Education Level	Marital Status	Income	Credit Score	Loan Amount	Loan Purpose	Employment Status	Years at Current Job	Payment History	Debt-to-Income Ratio	Assets Value	Number of Dependents	City
count	15000.000000	15000	15000	15000	12750.000000	12750.000000	12750.000000	15000	15000	15000.000000	15000	15000.000000	12750.000000	12750.000000	15000
unique	NaN	3	4	4	NaN	NaN	NaN	4	3	NaN	4	NaN	NaN	NaN	10614
top	NaN	Non-binary	Bachelor's	Widowed	NaN	NaN	NaN	Personal	Employed	NaN	Good	NaN	NaN	NaN	East Michael
freq	NaN	5059	3829	3893	NaN	NaN	NaN	3771	5026	NaN	3822	NaN	NaN	NaN	19
mean	43.452667	NaN	NaN	NaN	69933.398510	699.109098	27450.010902	NaN	NaN	9.476267	NaN	0.350438	159741.497176	2.02651	NaN
std	14.910732	NaN	NaN	NaN	29163.626207	57.229465	12949.940135	NaN	NaN	5.769707	NaN	0.143919	80298.115832	1.41130	NaN
min	18.000000	NaN	NaN	NaN	20005.000000	600.000000	5000.000000	NaN	NaN	0.000000	NaN	0.100004	20055.000000	0.000000	NaN
25%	31.000000	NaN	NaN	NaN	44281.500000	650.000000	16352.500000	NaN	NaN	4.000000	NaN	0.227386	90635.250000	1.000000	NaN
50%	43.000000	NaN	NaN	NaN	69773.000000	699.000000	27544.000000	NaN	NaN	9.000000	NaN	0.350754	159362.000000	2.000000	NaN
75%	56.000000	NaN	NaN	NaN	95922.750000	748.000000	38547.500000	NaN	NaN	15.000000	NaN	0.476095	228707.000000	3.000000	NaN
max	69.000000	NaN	NaN	NaN	119997.000000	799.000000	49998.000000	NaN	NaN	19.000000	NaN	0.599970	299999.000000	4.000000	NaN

5) Dropping the columns

dataset.drop() is used to remove specified rows or columns from the dataset.

- dataset.drop(columns=['column_name']) → Drops a specific column.
- dataset.drop(index=[row_index]) → Drops a specific row.

dropping the columns that aren't useful

```
cols = ['Marital Status', 'Marital Status Change', 'Loan Purpose', 'City', 'State']
df = dataset.drop(cols, axis=1)
df.info()
```

<class 'pandas.core.frame.DataFrame'>
 RangeIndex: 15000 entries, 0 to 14999
 Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	Age	15000 non-null	int64
1	Gender	15000 non-null	object
2	Education Level	15000 non-null	object
3	Income	12750 non-null	float64
4	Credit Score	12750 non-null	float64
5	Loan Amount	12750 non-null	float64
6	Employment Status	15000 non-null	object
7	Years at Current Job	15000 non-null	int64
8	Payment History	15000 non-null	object
9	Debt-to-Income Ratio	15000 non-null	float64
10	Assets Value	12750 non-null	float64
11	Number of Dependents	12750 non-null	float64
12	Country	15000 non-null	object
13	Previous Defaults	12750 non-null	float64
14	Risk Rating	15000 non-null	object

dtypes: float64(7), int64(2), object(6)
 memory usage: 1.7+ MB

Before Dropping:

```
[ ] dataset.shape
```

```
(15000, 20)
```

After Dropping:

```
df.shape
```

```
(15000, 15)
```

As observed here, the columns of 'Marital Status', 'Marital Status Change', 'Loan Purpose', 'City', 'State' have been dropped.

6) Drop rows with maximum missing rows

```
df["missing_count"] = df.isnull().sum(axis=1)
```

```
max_missing = df["missing_count"].max()
```

Here the maximum missing count is 6. So to clean up some of the data, we will remove the rows with 4 or more missing values. `df = df[df["missing_count"] < 4]`

The above set of commands do the following function:

- Create a column called `missing_count` where the sum of all the cells having null values is stored.
- The maximum value from this `missing_count` column is considered for deletion
- Finally, we update the dataset by keeping the rows which have missing values less than a particular value

```
df["missing_count"] = df.isnull().sum(axis=1)
max_missing = df["missing_count"].max()
print(df.head())

df = df[df["missing_count"] < 4]
df.shape
```

```
(14909, 16)
```

	Age	Gender	Education Level	Income	Credit Score	Loan Amount	\
0	49	Male	PhD	72799.0	688.0	45713.0	
1	57	Female	Bachelor's	NaN	690.0	33835.0	
2	21	Non-binary	Master's	55687.0	600.0	36623.0	
3	59	Male	Bachelor's	26508.0	622.0	26541.0	
4	25	Non-binary	Bachelor's	49427.0	766.0	36528.0	

	Employment Status	Years at Current Job	Payment History	\
0	Unemployed	19	Poor	
1	Employed	6	Fair	
2	Employed	8	Fair	
3	Unemployed	2	Excellent	
4	Unemployed	10	Fair	

	Debt-to-Income Ratio	Assets Value	Number of Dependents	Country	\
0	0.154313	120228.0	0.0	Cyprus	
1	0.148920	55849.0	0.0	Turkmenistan	
2	0.362398	180700.0	3.0	Luxembourg	
3	0.454964	157319.0	3.0	Uganda	
4	0.143242	287140.0	NaN	Namibia	

	Previous Defaults	Risk Rating	missing_count
0	2.0	Low	0
1	3.0	Medium	1
2	3.0	Medium	0
3	4.0	Medium	0
4	3.0	Low	1


To check the total missing values in each columns.

`df.isnull().sum()` is used to check for missing values (NaN) in a dataset. Here's how it works:

`df.isnull()` creates a DataFrame of the same shape as `df`, where each value is True if it's missing (NaN) and False otherwise.

`.sum()` then counts the number of True values (missing values) in each column.

7) Take care of the missing values



	0
Age	0
Gender	0
Education Level	0
Income	2189
Credit Score	2193
Loan Amount	2187
Employment Status	0
Years at Current Job	0
Payment History	0
Debt-to-Income Ratio	0
Assets Value	2196
Number of Dependents	2185
Country	0
Previous Defaults	2182
Risk Rating	0
missing_count	0

So, there are many missing values, hence performing the next step.

- To take care of the missing data that has not been removed, one of the 2 methods can be used: If the feature is of a numeric data type, we can use either mean, median or mode of the feature. If the data is normally distributed, use mean, if it is skewed, use median, and if many values are repeated, use mode.
- If the feature contains different categories, there are 2 ways. Either fill it with the mode of the column, or add a custom value such as “Data Unavailable”.

```
[ ] # handling the missing data
df.fillna({'Income':df['Income'].median()},inplace=True)
```

```
[ ] df.fillna({'Credit Score':df['Credit Score'].median()},inplace=True)
```

To check the columns with missing values, using the following command

`df[df.isnull().any(axis=1)]` filters and returns all rows that contain at least one missing (NaN) value.

```
missing_rows = df[df.isnull().any(axis=1)]
print(missing_rows)
```

Empty DataFrame
Columns: [Age, Gender, Education Level, Income, Credit Score, Loan Amount, Employment Status, Years at Current Job, Payment History, Debt-to-Income Ratio, Assets Value, Number of De
Index: []

8) Creating dummy variables

`pd.get_dummies(df, columns=categorical_columns, prefix=categorical_columns, drop_first=False)` is used to convert categorical variables into one-hot encoded format. This transformation helps machine learning models process categorical data.

Breaking Down the Code:

`pd.get_dummies(df, columns=categorical_columns, prefix=categorical_columns, drop_first=False)`

- Converts each categorical column into multiple binary (0/1) columns, representing unique categories.
- `prefix=categorical_columns` ensures that the new columns have meaningful names.
- `drop_first=False` keeps all categories (if True, it drops the first category to avoid multicollinearity).
- for col in categorical_columns: `df_dummies[col] = df[col]`

This restores the original categorical columns back into `df_dummies`, so the dataset now contains both original and encoded versions.

```

categorical_columns = ['Risk Rating', 'Gender', 'Employment Status', 'Payment History']

df_dummies = pd.get_dummies(df, columns=categorical_columns, prefix=categorical_columns, drop_first=False)

for col in categorical_columns:
    df_dummies[col] = df[col]

print(df_dummies.head())

```

	Age	Education Level	Income	Credit Score	Loan Amount	\
0	49	PhD	72799.0	688.0	45713.0	
1	57	Bachelor's	69773.0	690.0	33835.0	
2	21	Master's	55687.0	600.0	36623.0	
3	59	Bachelor's	26508.0	622.0	26541.0	
5	30	PhD	69773.0	717.0	15613.0	

	Years at Current Job	Debt-to-Income Ratio	Assets Value	\
0	19	0.154313	120228.000000	
1	6	0.148920	55849.000000	
2	8	0.362398	180700.000000	
3	2	0.454964	157319.000000	
5	5	0.295984	159741.497176	

	Number of Dependents	Country	...	Employment Status_Self-employed	\
0	0.0	Cyprus	...	False	
1	0.0	Turkmenistan	...	False	
2	3.0	Luxembourg	...	False	
3	3.0	Uganda	...	False	
5	4.0	Iceland	...	False	

	Employment Status_Unemployed	Payment History_Excellent	\
0	True	False	
1	False	False	
2	False	False	
3	True	True	
5	True	False	

	Payment History_Fair	Payment History_Good	Payment History_Poor	\
0	False	False	True	
1	True	False	False	
2	True	False	False	
3	False	False	False	
5	True	False	False	

	Risk Rating	Gender	Employment Status	Payment History
0	Low	Male	Unemployed	Poor
1	Medium	Female	Employed	Fair
2	Medium	Non-binary	Employed	Fair
3	Medium	Male	Unemployed	Excellent
5	Medium	Non-binary	Unemployed	Fair

9) Detecting Outlier data

Using IQR Value:

In this method, we find the IQR value for the column; which is the difference between $Q1 - 1.5 * IQR$ and $Q3 + 1.5 * IQR$. This is a standard that is followed, the factor 1.5 can be modified between 1 to 3 based on the requirement.

Command:

```
Q1 = df['Data_Value'].quantile(0.25)
```

```
Q3 = df['Data_Value'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
lower_bound = Q1 - 1.5 * IQR
```

```
upper_bound = Q3 + 1.5 * IQR
```

```
outliers = df[(df['Data_Value'] < lower_bound) | (df['Data_Value'] > upper_bound)]
```

This method gives the outliers and hence can be removed.

Using Manual method:

Checking for Outlier data in Excel using different value ranges.

And then using the preprocessed data.

```
df.to_csv('financial_risk_preprocessed.csv', index=False)
```

```
[ ] cleaned_df = pd.read_csv('/content/financial_risk_preprocessed WITH DEL.csv')
```

10) Standardization and Normalization of columns

- StandardScaler: Standardizes features by removing the mean and scaling to unit variance.
- MinMaxScaler: Normalizes features to a fixed range (0 to 1 by default).

Standardize Column:

Using formula:

```
mean_value = df["Data_Value"].mean()
```

```
std_value = df["Data_Value"].std()
```

```
df["Standardized_Data_Value"] = (df["Data_Value"] - mean_value) / std_value
```

Using Library:

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
df['Standardized Data Value Scaler'] = scaler.fit_transform(df[['Data_Value']])
```

Normalize column:

Method 1:

```
Formula min_val = df['Data_Value'].min()
```

```
max_val = df['Data_Value'].max()
```

```
df['Data_Value_Normalized'] = (df['Data_Value'] - min_val) / (max_val - min_val)
```

Method 2:

```
Scaler library from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
```

```
df['Normalized Data Value Scaler'] = scaler.fit_transform(df[['Data_Value']])
```

Here, the columns, Income, Credit Score, and Loan Amount are standardized and normalized



```
[ ] from sklearn.preprocessing import StandardScaler, MinMaxScaler

standard_scaler = StandardScaler()
min_max_scaler = MinMaxScaler()

cleaned_df['Income'] = standard_scaler.fit_transform(cleaned_df[['Income']])
cleaned_df['Credit Score'] = standard_scaler.fit_transform(cleaned_df[['Credit Score']])

cleaned_df['Loan Amount'] = min_max_scaler.fit_transform(cleaned_df[['Loan Amount']])

print(cleaned_df[['Income', 'Credit Score', 'Loan Amount']].head())
```



	Income	Credit Score	Loan Amount
0	-0.015390	-0.209478	0.000914
1	-0.017087	-0.172080	0.000677
2	-0.024984	-1.854955	0.000732
3	-0.041344	-1.443586	0.000531
4	-0.017087	0.332782	0.000312

Conclusion:

In this experiment, we used pandas and scikit learn to preprocess data, perform normalization and standardization to make the dataset clean and efficient. Firstly, the dataset in the form of csv file was imported into the Collab and then using `df.info()`, information related about the features (columns) of the dataset and the data type of each of these columns. Using `df.head()`, the dataset in the form of dataframe can be viewed in which the top 5 values can be displayed. Then the missing values were detected using `df.isnull()` and performing handling methods like dropping the rows with missing values, and replacing missing values with mean, median or mode, the missing values were handled.

After handling missing values and encoding categorical variables using one-hot encoding (creating dummy variables), the dataset was further refined by detecting and removing outliers using the Interquartile Range (IQR) method. Outliers were identified based on their deviation from the first (Q1) and third quartile (Q3) thresholds and manual processing on the dataset using Excel, ensuring that extreme values did not affect the model's performance.

StandardScaler was used to transform the columns, Income, Credit Score to standardized and normalized by centering them around a mean of zero with unit variance, making them suitable for models that assume normally distributed data. On the other hand, MinMaxScaler was applied to Loan Amount to a fixed range between 0 and 1.